

```
In [2]: import requests
from bs4 import BeautifulSoup
from datetime import datetime, timedelta
import time
import pandas as pd
import numpy as np
import random
from tqdm import tqdm
import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import TensorDataset, DataLoader
from sklearn.model_selection import train_test_split, KFold
import matplotlib.pyplot as plt
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import OneHotEncoder, StandardScaler
```

Section 1: Data Scraping & Wrangling Data

```
In [67]: def daterange(start_date, end_date):
    for n in range(int((end_date - start_date).days) + 1):
        yield start_date + timedelta(days=n)

team_name_mapping = {
    "Cleveland": "CLE",
    "Charlotte": "CHO",
    "LA Lakers": "LAL",
    "New Orleans": "NOP",
    "Brooklyn": "BRK",
    "Oklahoma City": "OKC",
    "Milwaukee": "MIL",
    "Phoenix": "PHO",
    "New York": "NYK",
    "Atlanta": "ATL",
    "Boston": "BOS",
    "Chicago": "CHI",
    "Dallas": "DAL",
    "Denver": "DEN",
    "Detroit": "DET",
    "Golden State": "GSW",
    "Houston": "HOU",
    "Indiana": "IND",
    "LA Clippers": "LAC",
    "Memphis": "MEM",
    "Miami": "MIA",
    "Minnesota": "MIN",
    "Orlando": "ORL",
    "Philadelphia": "PHI",
    "Portland": "POR",
    "Sacramento": "SAC",
    "San Antonio": "SAS",
    "Toronto": "TOR",
    "Utah": "UTA",
```

```

    "Washington": "WAS",
}

def rate_limit(requests, per_minute, last_request_time):
    if len(requests) >= per_minute:
        time_since_oldest_request = datetime.now() - requests[0]
        if time_since_oldest_request < timedelta(minutes=1):
            sleep_time = (timedelta(minutes=1) - time_since_oldest_request).total_seconds()
            print(f"Rate limit reached, sleeping for {sleep_time} seconds.")
            time.sleep(sleep_time)
        requests.pop(0)
    requests.append(datetime.now())

request_times = []

def scrape_basic_box_score_stats(url, team_abbrev):
    headers = {
        'User-Agent': 'Mozilla/5.0 (Macintosh; Intel Mac OS X 10.15; rv:86.0) Gecko/20100101 Firefox/86.0'
    }
    response = requests.get(url, headers=headers)
    rate_limit(request_times, 9, datetime.now())
    if response.status_code != 200:
        print(f"Failed to retrieve data from {url}")
        return {}

    soup = BeautifulSoup(response.content, 'html.parser')
    table_id = f'box-{team_abbrev}-game-basic'
    table = soup.find('table', id=table_id)

    desired_stats = {
        'FGA': None,
        'FG_pct': None,
        'FTA': None,
        'FT_pct': None,
        'TRB': None,
        'fg3a': None,
        'fg3_pct': None,
        'TOV': None,
        'stl': None,
        'blk': None,
        'ast': None
    }

    if table:
        totals_row = table.find('tfoot').find('tr')
        if totals_row:
            for stat in desired_stats.keys():
                data_cell = totals_row.find('td', {'data-stat': stat.lower()})
                if data_cell:
                    desired_stats[stat] = data_cell.text.strip()

    return desired_stats

def scrape_general_game_info(date):
    headers = {
        'User-Agent': 'Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/91.0.4472.164 Safari/537.36'
    }
    url = f"https://www.basketball-reference.com/boxscores/?month={date.month:02d}&day={date.day:02d}"
    rate_limit(request_times, 9, datetime.now())
    response = requests.get(url, headers=headers)

```

```

if response.status_code != 200:
    print(f"Request failed for {url} with status code {response.status_code}")
    return []

soup = BeautifulSoup(response.content, 'html.parser')
games_info = []

game_summaries = soup.find_all('div', class_='game_summary expanded nohover')
for game_summary in game_summaries:
    teams = game_summary.find_all('tr', class_=['winner', 'loser'])
    box_score_link_tag = game_summary.find('p', class_='links').find('a', text='Box Score')
    if len(teams) == 2 and box_score_link_tag:
        box_score_link = "https://www.basketball-reference.com" + box_score_link_tag.get('href')
        team1, team1_score = teams[0].find('a').text.strip(), teams[0].find('td').text.strip()
        team2, team2_score = teams[1].find('a').text.strip(), teams[1].find('td').text.strip()
        game_info = {
            'date': date.strftime('%Y-%m-%d'),
            'team1': team1,
            'team1_score': team1_score,
            'team2': team2,
            'team2_score': team2_score,
            'box_score_link': box_score_link
        }

        game_info['team1_stats'] = scrape_basic_box_score_stats(box_score_link, team1)
        game_info['team2_stats'] = scrape_basic_box_score_stats(box_score_link, team2)
        games_info.append(game_info)

return games_info

```

```

In [ ]: #date range
start_date = datetime(2023, 11, 1)
end_date = datetime(2024, 4, 5)
total_days = (end_date - start_date).days + 1

#collecting all games data
all_games_data = []
for single_date in tqdm(daterange(start_date, end_date), total=total_days, desc = "Scraping dates"):
    games_info = scrape_general_game_info(single_date)
    all_games_data.extend(games_info)

expanded_games_data = []

for game in all_games_data:
    #extracting basic game info
    game_info = {
        'date': game['date'],
        'team1': game['team1'],
        'team1_score': game['team1_score'],
        'team2': game['team2'],
        'team2_score': game['team2_score'],
    }

    #extracting team1 stats and prefix with 'team1_'
    team1_stats = {f'team1_{k}': v for k, v in game['team1_stats'].items()}
    game_info.update(team1_stats)

    #extracting team2 stats and prefix with 'team2_'
    team2_stats = {f'team2_{k}': v for k, v in game['team2_stats'].items()}
    game_info.update(team2_stats)

```

```
#adding the expanded game info to the new list
expanded_games_data.append(game_info)

#converting to a DataFrame for analysis
df_games = pd.DataFrame(expanded_games_data)

print("Dataset with General Game Info and Basic Box Score Stats:")
print(df_games)
```

```
In [ ]: #converted scraped data to csv to save it... so i don't spend another 5 hours data scr

df_games.to_csv('br.csv', index = False)
```

```
In [68]: br_df = pd.read_csv('br.csv')
print(br_df)
```

	date	team1	team1_score	team2	team2_score	\
0	11/1/2023	Washington	121	Atlanta	130	
1	11/1/2023	Indiana	104	Boston	155	
2	11/1/2023	Chicago	105	Dallas	114	
3	11/1/2023	Portland	110	Detroit	101	
4	11/1/2023	Sacramento	101	Golden State	102	
...	
1098	4/5/2024	Detroit	90	Memphis	108	
1099	4/5/2024	Toronto	117	Milwaukee	111	
1100	4/5/2024	San Antonio	111	New Orleans	109	
1101	4/5/2024	Minnesota	87	Phoenix	97	
1102	4/5/2024	Portland	108	Washington	102	

	team1_FGA	team1_FG_pct	team1_FTA	team1_FT_pct	team1_TRB	...	\
0	101	0.505	14	0.643	35	...	
1	101	0.455	12	0.583	31	...	
2	89	0.472	11	0.909	43	...	
3	81	0.469	33	0.818	44	...	
4	88	0.409	23	0.783	48	...	
...	
1098	79	0.405	24	0.750	38	...	
1099	79	0.456	36	0.861	47	...	
1100	83	0.506	13	0.769	34	...	
1101	85	0.388	14	0.786	44	...	
1102	94	0.436	24	0.708	55	...	

	team2_FG_pct	team2_FTA	team2_FT_pct	team2_TRB	team2_fg3a	\
0	0.500	32	0.906	57	32	
1	0.568	28	0.964	57	35	
2	0.457	28	0.714	45	48	
3	0.444	11	0.909	38	29	
4	0.481	15	0.867	36	31	
...	
1098	0.467	16	0.875	49	31	
1099	0.436	21	0.810	45	42	
1100	0.512	12	1.000	41	29	
1101	0.434	29	0.759	49	26	
1102	0.432	30	0.700	45	33	

	team2_fg3_pct	team2_TOV	team2_stl	team2_blk	team2_ast
0	0.281	20	8	3	26
1	0.571	11	5	2	27
2	0.417	13	7	2	23
3	0.379	18	12	4	24
4	0.355	17	5	6	32
...
1098	0.323	20	13	7	24
1099	0.286	13	8	6	21
1100	0.310	17	11	2	28
1101	0.346	18	12	3	29
1102	0.152	9	9	5	24

[1103 rows x 27 columns]

```
In [69]: unique_teams = pd.unique(br_df[['team1', 'team2']].values.ravel('K'))
#mapped each team to a unique integer
team_to_id = {team: i for i, team in enumerate(unique_teams)}
print(team_to_id)
br_df['Home_Team_ID'] = br_df['team2'].map(team_to_id) # home team is team 2
br_df['Away_Team_ID'] = br_df['team1'].map(team_to_id) # away team is team 1
```

```
#added win column  
br_df['HomeWin'] = (br_df['team2_score'] > br_df['team1_score']).astype(int) #astype c  
  
br_df['date'] = pd.to_datetime(br_df['date'])  
br_df['month'] = (br_df['date'].dt.month - 11) % 12  
print(br_df)
```

```
{'Washington': 0, 'Indiana': 1, 'Chicago': 2, 'Portland': 3, 'Sacramento': 4, 'Charlotte': 5, 'LA Clippers': 6, 'Brooklyn': 7, 'Denver': 8, 'Cleveland': 9, 'New Orleans': 10, 'Milwaukee': 11, 'Memphis': 12, 'Detroit': 13, 'Toronto': 14, 'San Antonio': 15, 'Orlando': 16, 'Dallas': 17, 'New York': 18, 'Golden State': 19, 'Boston': 20, 'Utah': 21, 'Atlanta': 22, 'LA Lakers': 23, 'Phoenix': 24, 'Miami': 25, 'Philadelphia': 26, 'Oklahoma City': 27, 'Minnesota': 28, 'Houston': 29}
```

	date	team1	team1_score	team2	team2_score	\
0	2023-11-01	Washington	121	Atlanta	130	
1	2023-11-01	Indiana	104	Boston	155	
2	2023-11-01	Chicago	105	Dallas	114	
3	2023-11-01	Portland	110	Detroit	101	
4	2023-11-01	Sacramento	101	Golden State	102	
...	
1098	2024-04-05	Detroit	90	Memphis	108	
1099	2024-04-05	Toronto	117	Milwaukee	111	
1100	2024-04-05	San Antonio	111	New Orleans	109	
1101	2024-04-05	Minnesota	87	Phoenix	97	
1102	2024-04-05	Portland	108	Washington	102	

	team1_FGA	team1_FG_pct	team1_FTA	team1_FT_pct	team1_TRB	...	\
0	101	0.505	14	0.643	35	...	
1	101	0.455	12	0.583	31	...	
2	89	0.472	11	0.909	43	...	
3	81	0.469	33	0.818	44	...	
4	88	0.409	23	0.783	48	...	
...	
1098	79	0.405	24	0.750	38	...	
1099	79	0.456	36	0.861	47	...	
1100	83	0.506	13	0.769	34	...	
1101	85	0.388	14	0.786	44	...	
1102	94	0.436	24	0.708	55	...	

	team2_fg3a	team2_fg3_pct	team2_TOV	team2_stl	team2_blk	team2_ast	\
0	32	0.281	20	8	3	26	
1	35	0.571	11	5	2	27	
2	48	0.417	13	7	2	23	
3	29	0.379	18	12	4	24	
4	31	0.355	17	5	6	32	
...	
1098	31	0.323	20	13	7	24	
1099	42	0.286	13	8	6	21	
1100	29	0.310	17	11	2	28	
1101	26	0.346	18	12	3	29	
1102	33	0.152	9	9	5	24	

	Home_Team_ID	Away_Team_ID	HomeWin	month
0	22	0	1	0
1	20	1	1	0
2	17	2	1	0
3	13	3	0	0
4	19	4	1	0
...
1098	12	13	1	5
1099	11	14	0	5
1100	10	15	0	5
1101	24	28	1	5
1102	0	3	0	5

[1103 rows x 31 columns]

```
In [ ]: #I saved the new 31 column version of dataframe
br_df.to_csv('br31.csv', index = False)
```

Section 2: Predicting Wins with BR31 Dataframe

```
In [70]: home_team = torch.tensor(br_df['Home_Team_ID'], dtype=torch.long)
away_team = torch.tensor(br_df['Away_Team_ID'], dtype=torch.long)
num_classes = 30
home_teams_ohe = F.one_hot(home_team, num_classes=num_classes)
away_teams_ohe = F.one_hot(away_team, num_classes=num_classes)
date_stuff = torch.tensor(br_df['month'], dtype=torch.long)

months_ohe = F.one_hot(date_stuff, num_classes = 6)

x_ohe_tensor = torch.cat((home_teams_ohe, away_teams_ohe), dim=1)

numeric_data = br_df.drop(['date', 'month', 'Home_Team_ID', 'team1_score', 'team2_score'])
scaler = StandardScaler()
x_norm = scaler.fit_transform(numeric_data)
x_norm_tensor = torch.tensor(x_norm, dtype=torch.float32)

X_almost = torch.cat((x_norm_tensor, x_ohe_tensor), dim=1)
X = torch.cat((X_almost, months_ohe), dim = 1)

print(X.shape)

y = torch.tensor(br_df['HomeWin'], dtype=torch.long)

torch.Size([1103, 88])
```

```
In [71]: class Trainer:

    def __init__(self, model, opt_method, learning_rate, batch_size, epoch, l2):
        self.model = model
        if opt_method == "adam":
            self.optimizer = torch.optim.Adam(model.parameters(), learning_rate, weight_decay=l2)
        else:
            raise NotImplementedError("This optimization is not supported")

        self.epoch = epoch
        self.batch_size = batch_size

    def train(self, X_train, y_train, X_val, y_val, early_stop=True, draw_curve=True):
        train_dataset = TensorDataset(X_train, y_train)
        train_loader = DataLoader(train_dataset, batch_size=self.batch_size, shuffle=True)
        val_dataset = TensorDataset(X_val, y_val)
        val_loader = DataLoader(val_dataset, batch_size=self.batch_size, shuffle=False)

        train_loss_list, train_acc_list = [], []
        val_loss_list, val_acc_list = [], []
        weights = self.model.state_dict()
        lowest_val_loss = np.inf
        loss_func = nn.CrossEntropyLoss()
        for n in tqdm(range(self.epoch), leave=True):
            # enable train mode
```



```

self.model.train()
epoch_loss, epoch_acc = 0.0, 0.0
for X_batch, y_batch in train_loader:
    batch_importance = y_batch.shape[0] / len(train_dataset)
    y_pred = self.model(X_batch)
    batch_loss = loss_func(y_pred, y_batch)

    self.optimizer.zero_grad()
    batch_loss.backward()
    self.optimizer.step()

    epoch_loss += batch_loss.detach().cpu().item() * batch_importance
    batch_acc = torch.sum(torch.argmax(y_pred, axis=1) == y_batch) / y_batch.shape[0]
    epoch_acc += batch_acc.detach().cpu().item() * batch_importance
train_loss_list.append(epoch_loss)
train_acc_list.append(epoch_acc)
val_loss, val_acc = self.evaluate(val_dataset)
val_loss_list.append(val_loss)
val_acc_list.append(val_acc)

if early_stop:
    if val_loss < lowest_val_loss:
        lowest_val_loss = val_loss
        weights = self.model.state_dict()

if draw_curve:
    x_axis = np.arange(self.epoch)
    fig, axes = plt.subplots(1, 2, figsize=(10, 4))
    axes[0].plot(x_axis, train_loss_list, label="Train")
    axes[0].plot(x_axis, val_loss_list, label="Validation")
    axes[0].set_title("Loss")
    axes[0].legend()
    axes[1].plot(x_axis, train_acc_list, label="Train")
    axes[1].plot(x_axis, val_acc_list, label="Validation")
    axes[1].set_title("Accuracy")
    axes[1].legend()
    print(f"Validation accuracy: {np.mean(val_acc_list)} +/- {np.std(val_acc_list)}")

if early_stop:
    self.model.load_state_dict(weights)

return {
    "train_loss_list": train_loss_list,
    "train_acc_list": train_acc_list,
    "val_loss_list": val_loss_list,
    "val_acc_list": val_acc_list,
}

def evaluate(self, data, print_acc=False):
    # enable evaluation mode
    self.model.eval()
    loader = DataLoader(data, batch_size=self.batch_size, shuffle=True)
    loss_func = nn.CrossEntropyLoss()
    acc, loss = 0.0, 0.0
    for X_batch, y_batch in loader:
        with torch.no_grad():
            batch_importance = y_batch.shape[0] / len(data)
            y_pred = self.model(X_batch)
            batch_loss = loss_func(y_pred, y_batch)
            batch_acc = torch.sum(torch.argmax(y_pred, axis=1) == y_batch) / y_batch.shape[0]

```

```

        acc += batch_acc.detach().cpu().item() * batch_importance
        loss += batch_loss.detach().cpu().item() * batch_importance
    if print_acc:
        print(f"Accuracy: {acc:.3f}")
    return loss, acc

```

```

In [72]: def KFoldCrossValidation(
    model_class, k,
    X, y,
    opt_method='adam', learning_rate=1e-4, batch_size=32, epoch=100, l2=0
):
    _, X_test, _, y_test = train_test_split(X, y, test_size=0.2)
    test_dataset = TensorDataset(X_test, y_test)
    kf = KFold(n_splits = k, shuffle = True)
    train_acc_list, test_acc_list = [], []
    for i, (train_index, val_index) in enumerate(kf.split(X)):
        print(f"Fold {i}:")

        X_train, X_val = X[train_index], X[val_index]
        y_train, y_val = y[train_index], y[val_index]

        model = model_class()
        # initialize a Trainer object
        trainer = Trainer(model, opt_method, learning_rate, batch_size, epoch, l2)
        # call trainer.train() here
        res = trainer.train(X_train, y_train, X_val, y_val)
        train_acc_best = res['train_acc_list'][np.argmin(res['val_loss_list'])]
        test_loss, test_acc = trainer.evaluate(test_dataset)
        if i == 1:
            torch.save(model.state_dict(), f'model_weights_fold_1.pth')
        train_acc_list.append(train_acc_best)
        test_acc_list.append(test_acc)

        print(f"Training accuracy: {train_acc_best}")
        print(f"Test accuracy: {test_acc}")

    print("Final results:")
    print(f"Training accuracy: {np.mean(train_acc_list)}+/-{np.std(train_acc_list)}")
    print(f"Test accuracy: {np.mean(test_acc_list)}+/-{np.std(test_acc_list)}")

```

```

In [73]: class WinPredictor(nn.Module):
    def __init__(self, ):
        super(WinPredictor, self).__init__()

        self.fc1 = nn.Linear(88, 128)

        self.fc2 = nn.Linear(128, 64)

        self.output_layer = nn.Linear(64, 2)

    def forward(self, x):

        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        # Using softmax activation function
        win_prob = F.softmax(self.output_layer(x), dim = -1)

        return win_prob

```

```
In [75]: KFoldCrossValidation(WinPredictor, 3, X, y)
```

Fold 0:

100%|██████████| 100/100 [00:14<00:00, 6.96it/s]

Validation accuracy: 0.9277445652173916+/-0.07012459269549368

Training accuracy: 1.0

Test accuracy: 0.9864253393665159

Fold 1:

100%|██████████| 100/100 [00:10<00:00, 9.62it/s]

Validation accuracy: 0.9235054347826086+/-0.05982628715643639

Training accuracy: 1.0

Test accuracy: 0.9773755631835214

Fold 2:

100%|██████████| 100/100 [00:10<00:00, 9.46it/s]

Validation accuracy: 0.8991553135691287+/-0.0688608152054051





Training accuracy: 0.9999999999999996


Test accuracy: 0.9819004500613495

Final results:

Training accuracy: 0.9999999999999999+/-2.1259152388627147e-16

Test accuracy: 0.9819004508704623+/-0.0036945556557881403





Analysis:

1. High variance/overfitting. Model is not learning underlying relationships.
2. Given team stats for every game, model predicts wins very well (i.e., much better than random chance).

Section 2: Creating rolling average and over/under prediction

```
In [76]: #prepared stats for team1 and team2
br_df = pd.read_csv('br31.csv')
team1_stats = br_df.filter(regex='team1').rename(columns=lambda x: x.replace('team1_',
team1_stats['team'] = br_df['team1']
team1_stats['date'] = br_df['date']

team2_stats = br_df.filter(regex='team2').rename(columns=lambda x: x.replace('team2_',
team2_stats['team'] = br_df['team2']
team2_stats['date'] = br_df['date']

#concatenate and calculate cumulative stats
all_stats = pd.concat([team1_stats, team2_stats])
all_stats_grouped = all_stats.groupby(['team', 'date']).sum(numeric_only=True).groupby
all_stats_grouped['games_played'] = all_stats.groupby(['team', 'date']).size().groupby

#ensured indices are aligned properly
current_game_stats = all_stats.groupby(['team', 'date']).sum(numeric_only=True)
current_game_stats = current_game_stats.reindex(all_stats_grouped.index)

#calculated running averages by excluding current game stats
all_stats_grouped['games_played'] = all_stats_grouped['games_played'] - 1 # Decrement
avg_stats_per_game = (all_stats_grouped - current_game_stats)
avg_stats_per_game = avg_stats_per_game.div(all_stats_grouped['games_played'], axis=0)
```

```
# Handle division by zero for the first game
avg_stats_per_game = avg_stats_per_game.replace([np.inf, -np.inf], np.nan).fillna(0)
avg_stats_per_game.reset_index(inplace=True)

print(avg_stats_per_game)
```

	team	date	FGA	FG_pct	FTA	FT_pct	\
0	Atlanta	2023-11-01	0.000000	0.000000	0.000000	0.000000	
1	Atlanta	2023-11-04	92.000000	0.500000	32.000000	0.906000	
2	Atlanta	2023-11-06	92.500000	0.492000	28.000000	0.849000	
3	Atlanta	2023-11-09	95.666667	0.452333	29.333333	0.847333	
4	Atlanta	2023-11-11	93.000000	0.459750	29.500000	0.827250	
...	
2201	Washington	2024-03-29	91.457143	0.471600	20.071429	0.762443	
2202	Washington	2024-03-31	91.408451	0.470718	19.901408	0.762268	
2203	Washington	2024-04-02	91.361111	0.470500	19.916667	0.762264	
2204	Washington	2024-04-03	91.383562	0.470973	19.876712	0.763110	
2205	Washington	2024-04-05	91.378378	0.471297	19.783784	0.765270	

	TOV	TRB	ast	blk	fg3_pct	fg3a	\
0	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
1	20.000000	57.000000	26.000000	3.000000	0.281000	32.000000	
2	15.500000	54.500000	27.000000	4.500000	0.311000	36.500000	
3	15.333333	56.000000	27.000000	4.000000	0.318333	38.333333	
4	15.750000	51.000000	25.750000	4.750000	0.335000	38.500000	
...	
2201	13.571429	40.885714	28.071429	5.100000	0.349043	35.157143	
2202	13.633803	40.971831	27.985915	5.084507	0.347972	35.126761	
2203	13.652778	40.986111	27.972222	5.027778	0.347153	35.166667	
2204	13.602740	41.027397	27.945205	4.986301	0.345644	35.205479	
2205	13.662162	41.000000	27.918919	4.972973	0.346770	35.297297	

	games_played	score	stl
0	0.0	0.000000	0.000000
1	0.0	130.000000	8.000000
2	0.0	126.500000	6.000000
3	0.0	123.333333	6.666667
4	0.0	122.500000	8.250000
...
2201	0.0	113.742857	7.571429
2202	0.0	113.366197	7.492958
2203	0.0	113.277778	7.472222
2204	0.0	113.328767	7.493151
2205	0.0	113.418919	7.540541

[2206 rows x 15 columns]

In [77]:

```
#Loading original data
br_df = pd.read_csv('br31.csv')

#making sure date is in datetime format
br_df['date'] = pd.to_datetime(br_df['date'])
avg_stats_per_game['date'] = pd.to_datetime(avg_stats_per_game['date'])

#creating unique indentifiers for merging
br_df['team_date'] = br_df['team1'].astype(str) + '_' + br_df['date'].dt.strftime('%Y-%m-%d')
avg_stats_per_game['team_date'] = avg_stats_per_game['team'].astype(str) + '_' + avg_stats_per_game['date'].dt.strftime('%Y-%m-%d')

#renaming date to avoid conflict
avg_stats_per_game = avg_stats_per_game.rename(columns={'date': 'date_avg'})
```

```
# Rename columns in avg_stats_per_game for team1
avg_stats_per_game_team1 = avg_stats_per_game.rename(columns=lambda x: x + '_avg1' if

# Merge the average statistics with the game results for team1
br_df = pd.merge(br_df, avg_stats_per_game_team1, on='team_date', how='left')

# Restore original 'date' column
br_df['date_avg1'] = br_df['date_avg']
br_df.drop(columns=['date_avg'], inplace=True)

# Update the unique identifiers for merging team2 statistics
br_df['team_date'] = br_df['team2'].astype(str) + '_' + br_df['date'].dt.strftime('%Y-

# Rename columns in avg_stats_per_game for team2
avg_stats_per_game_team2 = avg_stats_per_game.rename(columns=lambda x: x + '_avg2' if

# Merge the average statistics with the game results for team2
br_df = pd.merge(br_df, avg_stats_per_game_team2, on='team_date', how='left')

# Restore original 'date' column for team2
br_df['date_avg2'] = br_df['date_avg']
br_df.drop(columns=['date_avg'], inplace=True)

# Clean up unnecessary columns
columns_to_keep = [
    'date', 'team1', 'team1_score', 'team2', 'team2_score', 'Home_Team_ID', 'Away_Team_ID',
    'FGA_avg1', 'FG_pct_avg1', 'FTA_avg1', 'FT_pct_avg1', 'TOV_avg1', 'TRB_avg1', 'ast_avg1',
    'fg3_pct_avg1', 'fg3a_avg1', 'score_avg1', 'stl_avg1', 'FGA_avg2', 'FG_pct_avg2',
    'TOV_avg2', 'TRB_avg2', 'ast_avg2', 'blk_avg2', 'fg3_pct_avg2', 'fg3a_avg2', 'score_avg2'
]
br_df = br_df[columns_to_keep]

#saved the final DataFrame
#br_df.to_csv('br33.csv', index=False)

print(br_df)
```

	date	team1	team1_score	team2	team2_score	\
0	2023-11-01	Washington	121	Atlanta	130	
1	2023-11-01	Indiana	104	Boston	155	
2	2023-11-01	Chicago	105	Dallas	114	
3	2023-11-01	Portland	110	Detroit	101	
4	2023-11-01	Sacramento	101	Golden State	102	
...	
1098	2024-04-05	Detroit	90	Memphis	108	
1099	2024-04-05	Toronto	117	Milwaukee	111	
1100	2024-04-05	San Antonio	111	New Orleans	109	
1101	2024-04-05	Minnesota	87	Phoenix	97	
1102	2024-04-05	Portland	108	Washington	102	

	Home_Team_ID	Away_Team_ID	HomeWin	month	FGA_avg1	...	FTA_avg2	\
0	22	0	1	0	0.000000	...	0.000000	
1	20	1	1	0	0.000000	...	0.000000	
2	17	2	1	0	0.000000	...	0.000000	
3	13	3	0	0	0.000000	...	0.000000	
4	19	4	1	0	0.000000	...	0.000000	
...	
1098	12	13	1	5	88.333333	...	21.430556	
1099	11	14	0	5	89.875000	...	24.150685	
1100	10	15	0	5	90.847222	...	23.301370	
1101	24	28	1	5	84.863014	...	23.763889	
1102	0	3	0	5	89.236111	...	19.783784	

	FT_pct_avg2	TOV_avg2	TRB_avg2	ast_avg2	blk_avg2	fg3_pct_avg2	\
0	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
1	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
2	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
3	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
4	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
...	
1098	0.757750	14.180556	42.347222	25.041667	6.208333	0.346667	
1099	0.769904	11.986301	44.465753	26.890411	5.095890	0.375247	
1100	0.771411	12.260274	44.191781	27.136986	4.739726	0.377164	
1101	0.802083	13.930556	44.125000	27.236111	6.069444	0.382958	
1102	0.765270	13.662162	41.000000	27.918919	4.972973	0.346770	

	fg3a_avg2	score_avg2	stl_avg2
0	0.000000	0.000000	0.000000
1	0.000000	0.000000	0.000000
2	0.000000	0.000000	0.000000
3	0.000000	0.000000	0.000000
4	0.000000	0.000000	0.000000
...
1098	38.138889	106.013889	7.930556
1099	38.383562	120.136986	6.684932
1100	32.191781	115.589041	8.246575
1101	32.222222	117.194444	7.222222
1102	35.297297	113.418919	7.540541

[1103 rows x 33 columns]

Over/Under Prediction With RandomForestClassifier and BR33 Dataframe That Uses Rolling Averages

```
In [78]: br_df = pd.read_csv('br33.csv')
#dropping rows with 0s
br_df = br_df[(br_df['FGA_avg1'] != 0)
              ]
br_df = br_df.reset_index(drop = True)
threshold = 220
br_df['OverUnder'] = (br_df['team2_score'] + br_df['team1_score'] > threshold).astype(int)
#astype coverts boolean to integer 0 and 1

print(br_df)
br_cool = br_df
```


	date	team1	team1_score	team2	team2_score	\
0	2023-11-02	Detroit	116	New Orleans	125	
1	2023-11-02	Toronto	99	Philadelphia	114	
2	2023-11-03	Brooklyn	109	Chicago	107	
3	2023-11-03	Dallas	114	Denver	125	
4	2023-11-03	Cleveland	116	Indiana	121	
...	
1083	2024-04-05	Detroit	90	Memphis	108	
1084	2024-04-05	Toronto	117	Milwaukee	111	
1085	2024-04-05	San Antonio	111	New Orleans	109	
1086	2024-04-05	Minnesota	87	Phoenix	97	
1087	2024-04-05	Portland	108	Washington	102	

	Home_Team_ID	Away_Team_ID	HomeWin	month	FGA_avg1	...	FT_pct_avg2	\
0	10	13	1	0	90.000000	...	0.586000	
1	26	14	1	0	91.000000	...	0.000000	
2	2	7	0	0	82.000000	...	0.909000	
3	8	17	1	0	81.000000	...	0.700000	
4	1	9	1	0	74.000000	...	0.583000	
...	
1083	12	13	1	5	88.333333	...	0.757750	
1084	11	14	0	5	89.875000	...	0.769904	
1085	10	15	0	5	90.847222	...	0.771411	
1086	24	28	1	5	84.863014	...	0.802083	
1087	0	3	0	5	89.236111	...	0.765270	

	TOV_avg2	TRB_avg2	ast_avg2	blk_avg2	fg3_pct_avg2	fg3a_avg2	\
0	10.000000	58.000000	23.000000	5.000000	0.310000	42.000000	
1	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
2	13.000000	43.000000	19.000000	3.000000	0.333000	33.000000	
3	15.000000	43.000000	23.000000	10.000000	0.182000	33.000000	
4	9.000000	31.000000	26.000000	6.000000	0.135000	37.000000	
...	
1083	14.180556	42.347222	25.041667	6.208333	0.346667	38.138889	
1084	11.986301	44.465753	26.890411	5.095890	0.375247	38.383562	
1085	12.260274	44.191781	27.136986	4.739726	0.377164	32.191781	
1086	13.930556	44.125000	27.236111	6.069444	0.382958	32.222222	
1087	13.662162	41.000000	27.918919	4.972973	0.346770	35.297297	

	score_avg2	stl_avg2	OverUnder
0	110.000000	8.000000	1
1	0.000000	0.000000	0
2	105.000000	10.000000	0
3	89.000000	6.000000	1
4	104.000000	6.000000	1
...
1083	106.013889	7.930556	0
1084	120.136986	6.684932	1
1085	115.589041	8.246575	0
1086	117.194444	7.222222	0
1087	113.418919	7.540541	0

[1088 rows x 34 columns]

```
In [79]: numeric_data = br_df.drop(['date', 'month', 'team1_score', 'team2_score', 'OverUnder',
                                   'HomeWin', 'team1', 'team2'], axis = 1).values

X = numeric_data
```

```
y = br_df['OverUnder'].values

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=
```

```
In [80]: rf = RandomForestClassifier(random_state=42, n_jobs=-1)

params = {
    'max_depth': [2,3,5,10,20],
    'min_samples_leaf': [5,10,20,50,100,200],
    'n_estimators': [10,25,30,50,100,200]
}

from sklearn.model_selection import GridSearchCV

# Instantiate the grid search model
grid_search = GridSearchCV(estimator=rf,
                           param_grid=params,
                           cv = 4,
                           n_jobs=-1, verbose=1, scoring="accuracy")

grid_search.fit(X_train, y_train)
```

Fitting 4 folds for each of 180 candidates, totalling 720 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 46 tasks      | elapsed: 11.4s
[Parallel(n_jobs=-1)]: Done 196 tasks    | elapsed: 34.1s
[Parallel(n_jobs=-1)]: Done 446 tasks    | elapsed: 1.3min
[Parallel(n_jobs=-1)]: Done 720 out of 720 | elapsed: 2.2min finished
```

```
Out[80]: GridSearchCV(cv=4, error_score=nan,
                      estimator=RandomForestClassifier(bootstrap=True, ccp_alpha=0.0,
                                                         class_weight=None,
                                                         criterion='gini', max_depth=None,
                                                         max_features='auto',
                                                         max_leaf_nodes=None,
                                                         max_samples=None,
                                                         min_impurity_decrease=0.0,
                                                         min_impurity_split=None,
                                                         min_samples_leaf=1,
                                                         min_samples_split=2,
                                                         min_weight_fraction_leaf=0.0,
                                                         n_estimators=100, n_jobs=-1,
                                                         oob_score=False, random_state=42,
                                                         verbose=0, warm_start=False),
                      iid='deprecated', n_jobs=-1,
                      param_grid={'max_depth': [2, 3, 5, 10, 20],
                                   'min_samples_leaf': [5, 10, 20, 50, 100, 200],
                                   'n_estimators': [10, 25, 30, 50, 100, 200]},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                      scoring='accuracy', verbose=1)
```

```
In [81]: print(grid_search.best_score_)
          rf_best = grid_search.best_estimator_
          print(rf_best)
```

0.6459909947998139

```
RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                        criterion='gini', max_depth=5, max_features='auto',
                        max_leaf_nodes=None, max_samples=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=20, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=200,
                        n_jobs=-1, oob_score=False, random_state=42, verbose=0,
                        warm_start=False)
```

```
In [82]: # Instantiate the model
rf = RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                           criterion='gini', max_depth=5, max_features='auto',
                           max_leaf_nodes=None, max_samples=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=20, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=200,
                           n_jobs=-1, oob_score=False, verbose=0,
                           warm_start=False)

# Train the model
rf.fit(X_train, y_train)

# Make predictions
y_pred = rf.predict(X_test)
#print(rf.oob_score_)
# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy}')
```

Accuracy: 0.6697247706422018

With standard scaling and onehot encoding

```
In [83]: cat = br_df[['Home_Team_ID', 'Away_Team_ID']].values
encoder = OneHotEncoder()
cat_ohe = encoder.fit_transform(cat).toarray()

numeric_data = br_df.drop(['date', 'month', 'team1_score', 'team2_score', 'OverUnder',
                           'Home_Team_ID', 'Away_Team_ID', 'HomeWin', 'team1', 'team2'])

scaler = StandardScaler()
x_norm = scaler.fit_transform(numeric_data)
X = np.hstack((x_norm, cat_ohe))
print(X.shape)

y = br_df['OverUnder'].values

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Instantiate the model
rf = RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                           criterion='gini', max_depth=5, max_features='auto',
                           max_leaf_nodes=None, max_samples=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=20, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=200,
                           n_jobs=-1, oob_score=False, verbose=0,
```

```
warm_start=False)  
  
# Train the model  
rf.fit(X_train, y_train)  
  
# Make predictions  
y_pred = rf.predict(X_test)  
#print(rf.oob_score_)  
# Evaluate the model  
accuracy = accuracy_score(y_test, y_pred)  
print(f'Accuracy: {accuracy}')
```

(1088, 84)

Accuracy: 0.6880733944954128

Individual match prediction

```
In [84]: br_df = pd.read_csv('br33.csv')  
#dropping rows with 0s  
br_df = br_df[(br_df['FGA_avg1'] != 0)  
              ]  
br_df = br_df.reset_index(drop = True)  
  
print(br_df)  
br_cool = br_df
```

	date	team1	team1_score	team2	team2_score	\
0	2023-11-02	Detroit	116	New Orleans	125	
1	2023-11-02	Toronto	99	Philadelphia	114	
2	2023-11-03	Brooklyn	109	Chicago	107	
3	2023-11-03	Dallas	114	Denver	125	
4	2023-11-03	Cleveland	116	Indiana	121	
...	
1083	2024-04-05	Detroit	90	Memphis	108	
1084	2024-04-05	Toronto	117	Milwaukee	111	
1085	2024-04-05	San Antonio	111	New Orleans	109	
1086	2024-04-05	Minnesota	87	Phoenix	97	
1087	2024-04-05	Portland	108	Washington	102	

	Home_Team_ID	Away_Team_ID	HomeWin	month	FGA_avg1	...	FTA_avg2	\
0	10	13	1	0	90.000000	...	29.000000	
1	26	14	1	0	91.000000	...	0.000000	
2	2	7	0	0	82.000000	...	11.000000	
3	8	17	1	0	81.000000	...	10.000000	
4	1	9	1	0	74.000000	...	12.000000	
...	
1083	12	13	1	5	88.333333	...	21.430556	
1084	11	14	0	5	89.875000	...	24.150685	
1085	10	15	0	5	90.847222	...	23.301370	
1086	24	28	1	5	84.863014	...	23.763889	
1087	0	3	0	5	89.236111	...	19.783784	

	FT_pct_avg2	TOV_avg2	TRB_avg2	ast_avg2	blk_avg2	fg3_pct_avg2	\
0	0.586000	10.000000	58.000000	23.000000	5.000000	0.310000	
1	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
2	0.909000	13.000000	43.000000	19.000000	3.000000	0.333000	
3	0.700000	15.000000	43.000000	23.000000	10.000000	0.182000	
4	0.583000	9.000000	31.000000	26.000000	6.000000	0.135000	
...	
1083	0.757750	14.180556	42.347222	25.041667	6.208333	0.346667	
1084	0.769904	11.986301	44.465753	26.890411	5.095890	0.375247	
1085	0.771411	12.260274	44.191781	27.136986	4.739726	0.377164	
1086	0.802083	13.930556	44.125000	27.236111	6.069444	0.382958	
1087	0.765270	13.662162	41.000000	27.918919	4.972973	0.346770	

	fg3a_avg2	score_avg2	stl_avg2
0	42.000000	110.000000	8.000000
1	0.000000	0.000000	0.000000
2	33.000000	105.000000	10.000000
3	33.000000	89.000000	6.000000
4	37.000000	104.000000	6.000000
...
1083	38.138889	106.013889	7.930556
1084	38.383562	120.136986	6.684932
1085	32.191781	115.589041	8.246575
1086	32.222222	117.194444	7.222222
1087	35.297297	113.418919	7.540541

[1088 rows x 33 columns]

With OHE and Standard Scaling

```
In [59]: team_names = {
          0: 'Washington',
```

```

1: 'Indiana',
2: 'Chicago',
3: 'Portland',
4: 'Sacramento',
5: 'Charlotte',
6: 'LA Clippers',
7: 'Brooklyn',
8: 'Denver',
9: 'Cleveland',
10: 'New Orleans',
11: 'Milwaukee',
12: 'Memphis',
13: 'Detroit',
14: 'Toronto',
15: 'San Antonio',
16: 'Orlando',
17: 'Dallas',
18: 'New York',
19: 'Golden State',
20: 'Boston',
21: 'Utah',
22: 'Atlanta',
23: 'LA Lakers',
24: 'Phoenix',
25: 'Miami',
26: 'Philadelphia',
27: 'Oklahoma City',
28: 'Minnesota',
29: 'Houston'
}

def prepare_features(home_stats, away_stats):
    """
    Prepares the numeric feature vector from home and away team stats.
    Uses suffixes '_1' for home team stats and '_2' for away team stats as per the dat
    """
    #stats list for home and away teams as per suffix conventions
    home_stats_list = [
        'FGA_avg1', 'FG_pct_avg1', 'FTA_avg1', 'FT_pct_avg1', 'TOV_avg1',
        'TRB_avg1', 'ast_avg1', 'blk_avg1', 'fg3_pct_avg1', 'fg3a_avg1',
        'score_avg1', 'stl_avg1'
    ]
    away_stats_list = [
        'FGA_avg2', 'FG_pct_avg2', 'FTA_avg2', 'FT_pct_avg2', 'TOV_avg2',
        'TRB_avg2', 'ast_avg2', 'blk_avg2', 'fg3_pct_avg2', 'fg3a_avg2',
        'score_avg2', 'stl_avg2'
    ]

    #extract the respective stats for home and away teams
    home_features = home_stats[home_stats_list].values.flatten()
    away_features = away_stats[away_stats_list].values.flatten()

    #combine home and away features
    return np.hstack([home_features, away_features])

def get_latest_team_stats(team_stats_df, team_id, home = None):
    """
    Retrieves the latest stats for a given team_id based on the most recent date.
    """
    if home is not None:

```

```

team_data = team_stats_df[team_stats_df['Home_Team_ID'] == team_id]
latest_entry = team_data.sort_values(by='date', ascending=False).iloc[0]
else:
    team_data = team_stats_df[team_stats_df['Away_Team_ID'] == team_id]
    latest_entry = team_data.sort_values(by='date', ascending=False).iloc[0]
return latest_entry

def predict_tomorrow_games(matchups, team_stats_df, thresholds):
    predictions = []

    all_team_ids = np.unique(team_stats_df[['Home_Team_ID', 'Away_Team_ID']].values.flatten())
    encoder = OneHotEncoder(categories=[all_team_ids, all_team_ids])
    scaler = StandardScaler()

    #preparing the numeric data for scaling and categorical data for encoding
    numeric_columns = team_stats_df.drop(['date',
                                           'month',
                                           'team1_score',
                                           'team2_score',
                                           'Home_Team_ID',
                                           'Away_Team_ID',
                                           'HomeWin',
                                           'team1',
                                           'team2'], axis = 1).values

    x_norm = scaler.fit_transform(numeric_columns)
    cat = team_stats_df[['Home_Team_ID', 'Away_Team_ID']].values
    cat_ohe = encoder.fit_transform(cat).toarray()
    X = np.hstack((x_norm, cat_ohe)) # Full feature set for training

    for index, (home_id, away_id) in enumerate(matchups):
        threshold = thresholds[index]
        probabilities_list = []
        accuracy_list = []
        team_stats_df2 = team_stats_df.copy(deep=True)

        for _ in range(20):

            team_stats_df2['OverUnder'] = (team_stats_df2['team2_score'] + team_stats_c
            y = team_stats_df2['OverUnder'].values

            #split data for training and testing
            X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25)
            rf = RandomForestClassifier(n_estimators=200)
            rf.fit(X_train, y_train)
            y_pred = rf.predict(X_test)
            accuracy = accuracy_score(y_test, y_pred)
            accuracy_list.append(accuracy)

            # prepare features for prediction for the current game
            home_stats = get_latest_team_stats(team_stats_df2, home_id, home=True)
            away_stats = get_latest_team_stats(team_stats_df2, away_id, home=None)
            numeric_features = prepare_features(home_stats, away_stats)
            numeric_features = scaler.fit_transform(numeric_features.reshape(1, -1))
            categorical_features = np.array([[home_id, away_id]])
            encoded_features = encoder.transform(categorical_features).toarray()

            game_features = np.hstack((encoded_features, numeric_features))

```

```

# Predict probabilities and store them
probabilities = rf.predict_proba(game_features)[0, 1]
probabilities_list.append(probabilities)

#average the probabilities and accuracies for current matchup
avg_probability = np.mean(probabilities_list)
avg_accuracy = np.mean(accuracy_list)
final_decision = "Over" if avg_probability > 0.5 else "Under"
confidence = "High Confidence" if avg_probability >= 0.60 or avg_probability < 0.40 else "Low Confidence"

predictions.append((team_names[home_id], team_names[away_id], round(avg_accuracy, 2), final_decision, confidence))

return predictions

```

```

In [60]: #march 20 to 21st ish 2023
# 12 correct and only 1 wrong
#WOW!

matchup = [(25, 2), (10, 4), (28, 2)] # Home team is (HERE, NOT HERE)
threshold = [207, 210, 200]

predict_tomorrow_games(matchups=matchup,
                        team_stats_df=br_cool,
                        thresholds=threshold)

```

```

Out[60]: [('Miami', 'Chicago', 0.858, 207, 0.944, 'Over', 'High Confidence'),
          ('New Orleans', 'Sacramento', 0.807, 210, 0.824, 'Over', 'High Confidence'),
          ('Minnesota', 'Chicago', 0.932, 200, 0.986, 'Over', 'High Confidence')]

```

Without Standard Scaling and One Hot Encoding

```

In [57]: team_names = {
    0: 'Washington', 1: 'Indiana', 2: 'Chicago', 3: 'Portland', 4: 'Sacramento',
    5: 'Charlotte', 6: 'LA Clippers', 7: 'Brooklyn', 8: 'Denver', 9: 'Cleveland',
    10: 'New Orleans', 11: 'Milwaukee', 12: 'Memphis', 13: 'Detroit', 14: 'Toronto',
    15: 'San Antonio', 16: 'Orlando', 17: 'Dallas', 18: 'New York', 19: 'Golden State',
    20: 'Boston', 21: 'Utah', 22: 'Atlanta', 23: 'LA Lakers', 24: 'Phoenix', 25: 'Miami',
    26: 'Philadelphia', 27: 'Oklahoma City', 28: 'Minnesota', 29: 'Houston'
}

def prepare_features(home_stats, away_stats):
    """
    Prepares the numeric feature vector from home and away team stats.
    Uses suffixes '_avg1' for home team stats and '_avg2' for away team stats as per the
    """

    home_stats_list = [
        'FGA_avg1', 'FG_pct_avg1', 'FTA_avg1', 'FT_pct_avg1', 'TOV_avg1',
        'TRB_avg1', 'ast_avg1', 'blk_avg1', 'fg3_pct_avg1', 'fg3a_avg1',
        'score_avg1', 'stl_avg1'
    ]

    away_stats_list = [
        'FGA_avg2', 'FG_pct_avg2', 'FTA_avg2', 'FT_pct_avg2', 'TOV_avg2',
        'TRB_avg2', 'ast_avg2', 'blk_avg2', 'fg3_pct_avg2', 'fg3a_avg2',
        'score_avg2', 'stl_avg2'
    ]

```



```

]

home_features = home_stats[home_stats_list].values.flatten()
away_features = away_stats[away_stats_list].values.flatten()

return np.hstack([home_features, away_features])

def get_latest_team_stats(team_stats_df, team_id, home=None):
    """
    Retrieves the latest stats for a given team_id based on the most recent date.
    """
    if home:
        team_data = team_stats_df[team_stats_df['Home_Team_ID'] == team_id]
    else:
        team_data = team_stats_df[team_stats_df['Away_Team_ID'] == team_id]

    latest_entry = team_data.sort_values(by='date', ascending=False).iloc[0]
    return latest_entry

def predict_tomorrow_games(matchups, team_stats_df, thresholds):
    predictions = []

    for index, (home_id, away_id) in enumerate(matchups):
        threshold = thresholds[index]
        probabilities_list = []
        accuracy_list = []

        team_stats_df2 = team_stats_df.copy(deep=True)
        for _ in range(20):
            team_stats_df2['OverUnder'] = (team_stats_df2['team2_score'] + team_stats_
            y = team_stats_df2['OverUnder'].values

            # Prepare numeric features and categorical team IDs for training
            numeric_columns = team_stats_df2.drop(columns=['date', 'month', 'team1_sco
            X_numeric = numeric_columns.values
            X_team_ids = team_stats_df2[['Home_Team_ID', 'Away_Team_ID']].values
            X = np.hstack((X_numeric, X_team_ids))

            X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
            rf = RandomForestClassifier(n_estimators=200, random_state=42)
            rf.fit(X_train, y_train)
            y_pred = rf.predict(X_test)
            accuracy = accuracy_score(y_test, y_pred)
            accuracy_list.append(accuracy)

            home_stats = get_latest_team_stats(team_stats_df2, home_id, home=True)
            away_stats = get_latest_team_stats(team_stats_df2, away_id, home=False)
            numeric_features = prepare_features(home_stats, away_stats)
            game_features = np.hstack((numeric_features, [home_id, away_id])).reshape(

            probabilities = rf.predict_proba(game_features)[0, 1]
            probabilities_list.append(probabilities)

        avg_probability = np.mean(probabilities_list)

```

```
avg_accuracy = np.mean(accuracy_list)
final_decision = "Over" if avg_probability > 0.5 else "Under"
confidence = "High Confidence" if avg_probability >= 0.60 or avg_probability < 0.40 else "Low Confidence"

predictions.append((team_names[home_id], team_names[away_id], round(avg_accuracy, 2), threshold, avg_probability, final_decision, confidence))

return predictions
```

```
In [58]: matchup = [(25, 2), (10, 4), (28, 2)] # Home team is (HERE, NOT HERE)
threshold = [207, 210, 200]

predict_tomorrow_games(matchups=matchup,
                        team_stats_df=br_cool,
                        thresholds=threshold)

Out[58]: [('Miami', 'Chicago', 0.86, 207, 0.815, 'Over', 'High Confidence'),
          ('New Orleans', 'Sacramento', 0.812, 210, 0.65, 'Over', 'High Confidence'),
          ('Minnesota', 'Chicago', 0.934, 200, 0.895, 'Over', 'High Confidence')]
```